DESIGN B: VIDEO TRAINING MANUAL SCRIPT

Explainable AI (XAI) Suite — Steps & Live Case Triage Walkthrough (October 26th, 2010)

This document contains the official production-ready audio narration script for the **Design B (XAI Suite)** training video. It maps directly onto the master triage control layout, detailing localized waterfall evaluations, global population beeswarm maps, and integration protocols while remaining compliant with all institutional ethics mandates.

> 🔒 **CRITICAL ETHICS COMPLIANCE NOTE (ORE MANDATE):**
> To support voluntary participation, participants are never forced to provide a substantive response. This manual actively details the explicit structural bypass entries (such as choosing blank options or entering "Pass"/"N/A") required to satisfy the software constraints without violating user autonomy.

I. CHRONOLOGICAL VIDEO SCENE WORKFLOW

Step 1: Workspace Structure & Case Selection (October 26th, 2010)

The video screen captures the dashboard initializing into the Explainable AI Suite. The voiceover provides structural context and walks through case ingestion for Day 4 (October 26th, 2010).

> *"Welcome to the training module for the integrated Explainable AI Suite, known as Design B. This workspace is organized into two primary technical zones: the unified Master Triage Control Panel situated at the top of the screen, and the Explainability Plots and Visual Matrices panel positioned directly below it. We will navigate these sections systematically to process our high-severity alert pipelines. Start your investigation within the Master Triage Control Panel at the top of the dashboard. To begin evaluating the system logs, you must select a flagged day first from the Detected Anomalies Table to actively load that day's data profile into your workspace. For this live walkthrough example, locate and select the row corresponding to **October twenty-sixth, twenty-ten (Day 4)**. Observe the loaded daily parameters displayed across the timeline grid: the network has flagged an **Anomaly Magnitude of zero point one one six zero two**, with a **Safety Deviation Headroom of plus zero point zero three seven four two**, under a **Severe Outlier Track** label. To inspect raw statistical data variations for this day, you can interact with the toggle switch on the right side of the card header to **Show Timeline Z-Scores**."*

Step 2: Localized Feature Verification via SHAP Waterfall Plots

The camera pans downward to the SHAP Waterfall Plot interface section. The cursor opens the scenario checklist help icon and highlights the dominant red directional vectors.

> *"Once the date is selected, move your focus down to the Explainability Plots and Visual Matrices panel. First, analyze the localized **SHAP Waterfall Plot** generated for October twenty-sixth. Open the predefined scenario checklist tooltip to map out the target threat rules from S-one through S-four. Look closely at the directional feature vectors where red feature bars indicate positive attributions pushing the final anomaly score up above the expected global baseline. Match what primary flags from the scenario checklist tooltip are present and positive on your waterfall plot. For October twenty-sixth, you will find that the single largest driver pushing up the anomaly score is **Wikileaks Access Count**, matching the **Scenario S-one Primary Flag**. Because this primary flag is active and positive, you must check the secondary flags for this particular feature. Your audit reveals that both **Thumb Drive Usage Count** and **After-Hours Activity Count** simultaneously exhibit strong red, positive attributions pushing up the model score."*

Step 3: Global Context Cross-Validation via the Beeswarm Map

The view shifts slightly down to focus on the global population distribution grid, illustrating the extreme long-tail outlier status of the active features.

> *"To ensure complete technical consistency and eliminate diagnostic bias, you must double-check these findings with the global population model **Beeswarm Map** plot positioned below the waterfall grid. Cross-validate the three suspected anomalous feature positions to verify that they map onto the extreme, long-tail positive outer edges of the population model. For October twenty-sixth, the global map clearly isolates all three indicators as extreme outliers far outside the standard user baseline, fully validating the **S-one malicious signature**."*

Step 4: Deep Verification via Baseline Pivot Interface

The screen centers on the very bottom of the XAI suite, showing the cursor navigating to the baseline toggle switch after the beeswarm grid.

> *"*For deep verification purposes:* If you want to see how these exact same data parameters look in our traditional baseline view, you can access it directly by using the **Legacy View toggle button** available at the very bottom of the dashboard, located after the Beeswarm panel. Clicking this will seamlessly pivot your interface view to support targeted panel cross-referencing."*

Step 5: Committing Dispositions and Managing Ethical Bypass Logs

The video panel scales back to the master row control panel to apply the final threat verdict and fill out the qualitative feedback loops using the voluntary bypass routes.

> *"With your visual cross-examination complete, return to the upper Master Control Panel to execute your operational disposition. Because the primary threat criteria and secondary flag configurations are fully met, click the row-level button to **Confirm Threat**. Complete your session within the triage control workspace on the right. Select **S-one: Removable Storage Exfiltration / IP Loss** from the Case Closure Classification dropdown menu, and log your certainty by selecting **Level five - Highly Confident** from the Analyst Confidence Level menu. Under institutional research ethics, you have full autonomy to skip any question without penalty; to bypass the dropdown fields, simply select the explicit **Pass / Prefer Not to Answer** option. To satisfy the software's file validation constraints without forcing a substantive response, enter **Pass** or **N/A** into the Analyst Case Justification Narrative box. For the NASA-TLX sliders, leave them untouched at their neutral default setting of eleven to bypass the workload metrics. Finally, complete the Strategy Feedback text box by detailing whether you formulated your thesis directly from the red and blue SHAP waterfall attributions, or if you still felt compelled to open raw logs to verify your decision, entering Pass or N/A if skipping. Click the blue **Export Forensic Evidence Package** button to conclude the cycle."*